

Goka

Painless stream processing with Go and Kafka

Franz Eichhorn Diogo Behrens

2019-01-29


Agenda

- Streaming microservices
 - problems and approach
- Mini Kafka review
 - basic concepts and hands-on warmup
- Learning Goka the “hard way”
 - components introduction with hands-on assignments
- Closing discussion
 - best practices, monitoring, testing, links, . . .


Streaming microservices

queues, DBs, logs, caches, . . .

Writing data: response time and complexity issues




Writing data: response time and complexity issues




Writing data: response time and complexity issues


Problem 1:
if DB slow, client slow




Writing data: response time and complexity issues




Writing data: response time and complexity issues



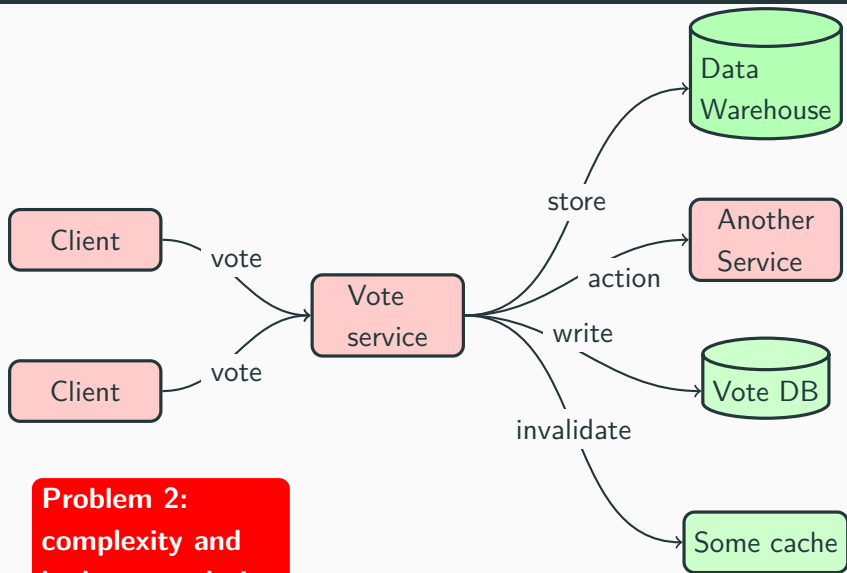
Writing data: response time and complexity issues



Writing data: response time and complexity issues

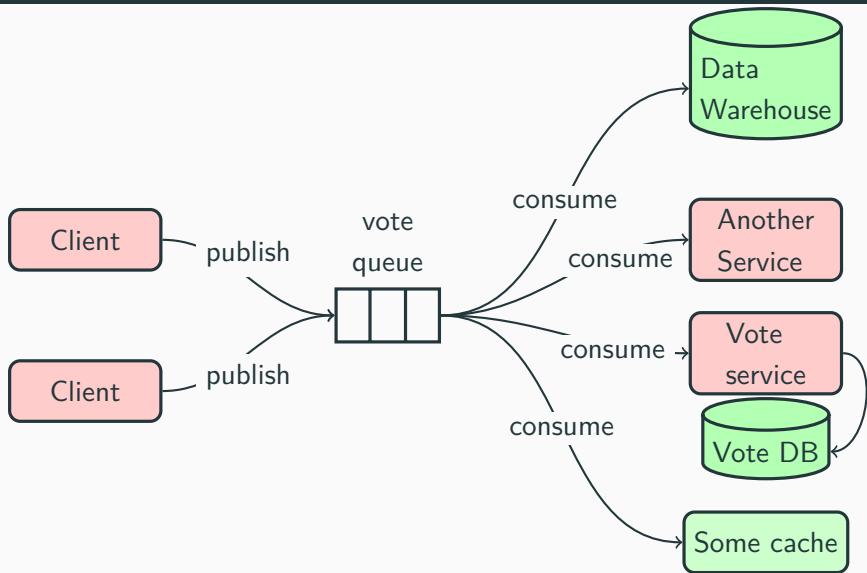


Writing data: response time and complexity issues



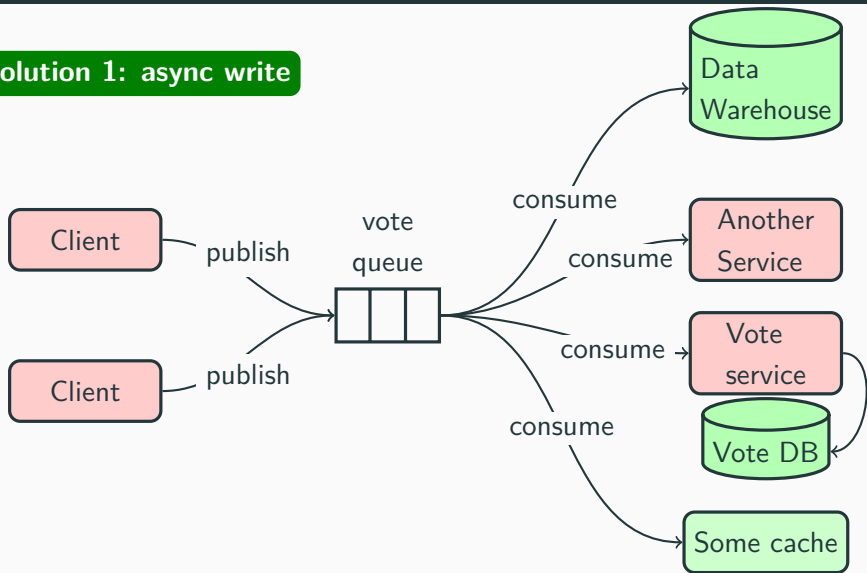
Problem 2:
complexity and
bad encapsulation

Writing data: decouple clients and services in space and time



Writing data: decouple clients and services in space and time

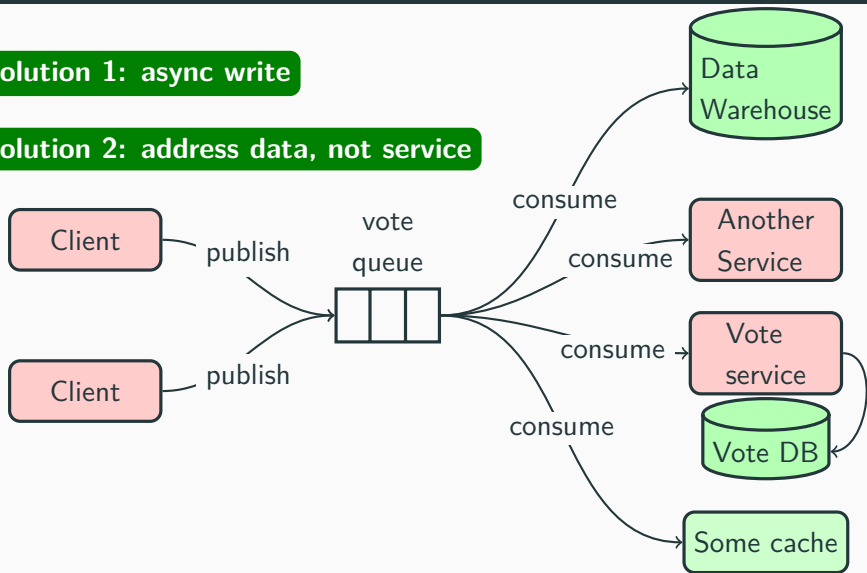
Solution 1: async write



Writing data: decouple clients and services in space and time

Solution 1: async write

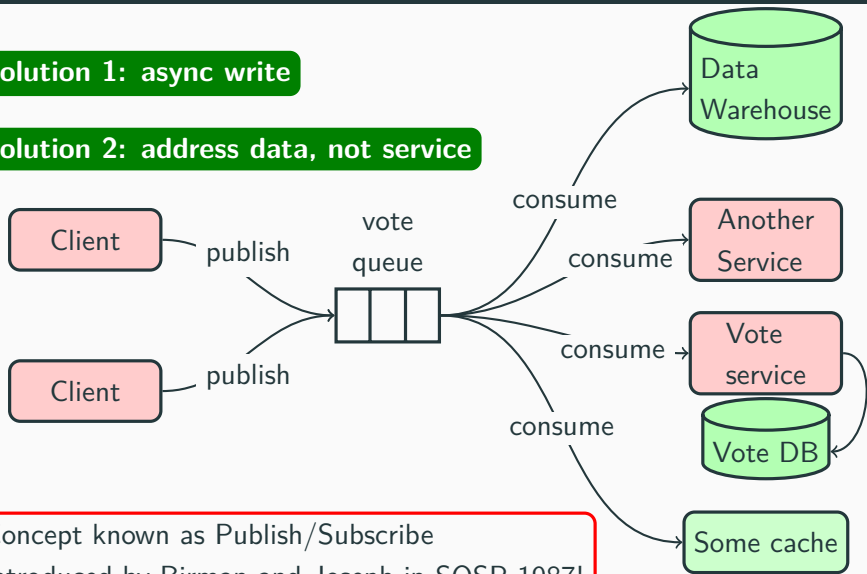
Solution 2: address data, not service



Writing data: decouple clients and services in space and time


Solution 1: async write

Solution 2: address data, not service




Concept known as Publish/Subscribe
Introduced by Birman and Joseph in SOSP 1987!


Reading data: microservice dependencies




Reading data: microservice dependencies




Reading data: microservice dependencies




Reading data: microservice dependencies



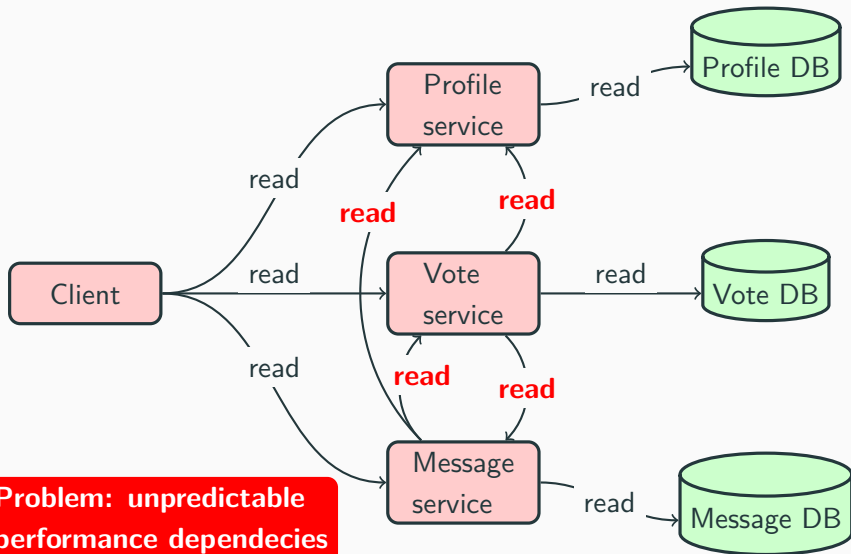
Reading data: microservice dependencies




Reading data: microservice dependencies




Reading data: microservice dependencies




Reading data: “old school” approach



Reading data: “old school” approach




Reading data: “old school” approach



Reading data: “old school” approach

Problem:


if DB slow, client slow



Reading data: “old school” approach

Solution:


if DB slow, add cache



Reading data: “old school” approach

Solution:

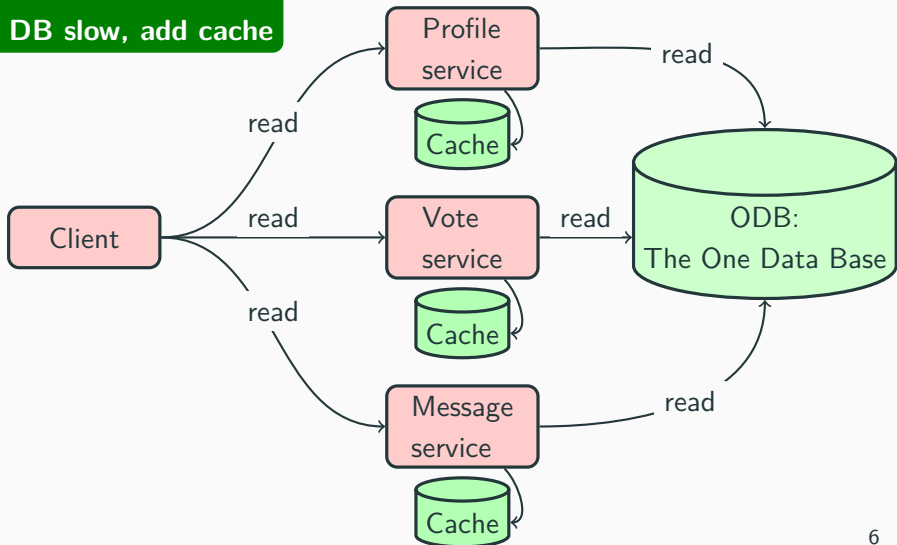
if DB slow, add cache



Reading data: “old school” approach

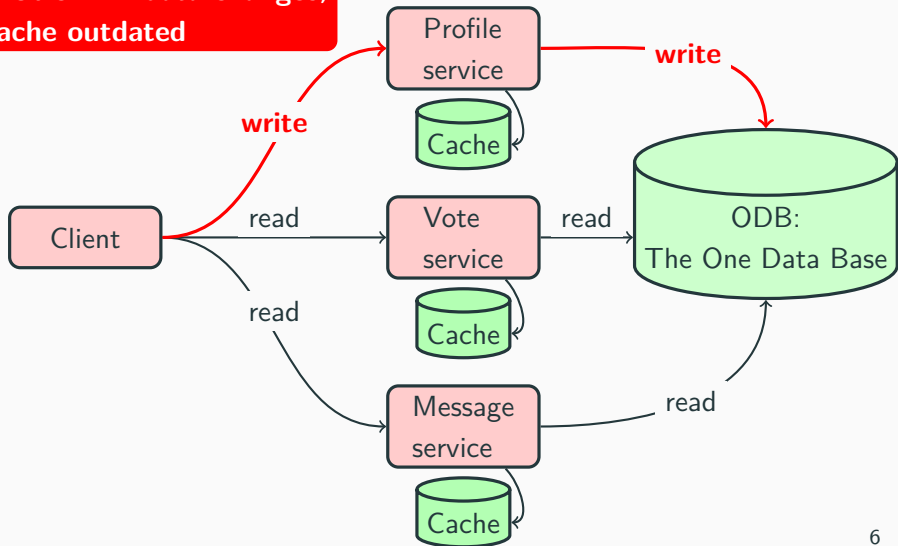
Solution:

if DB slow, add cache



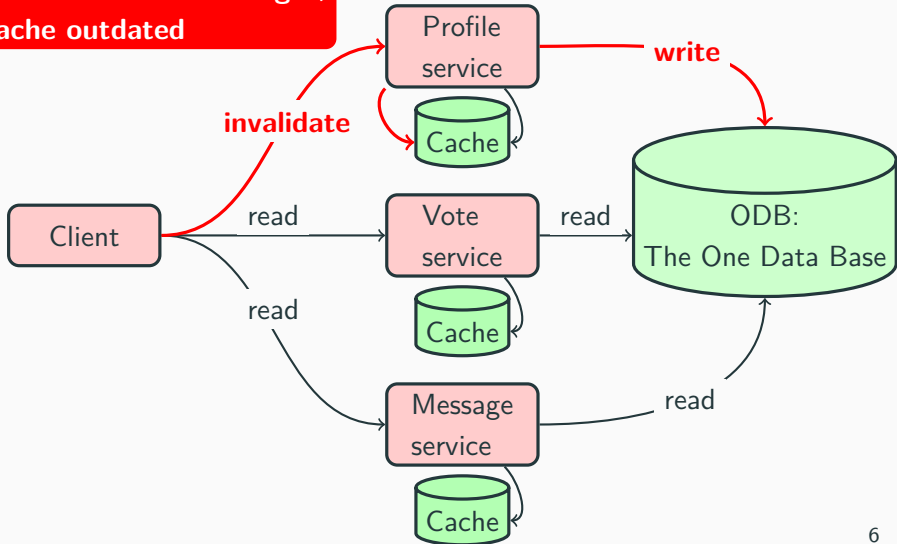
Reading data: “old school” approach

Problem: if data changes, cache outdated



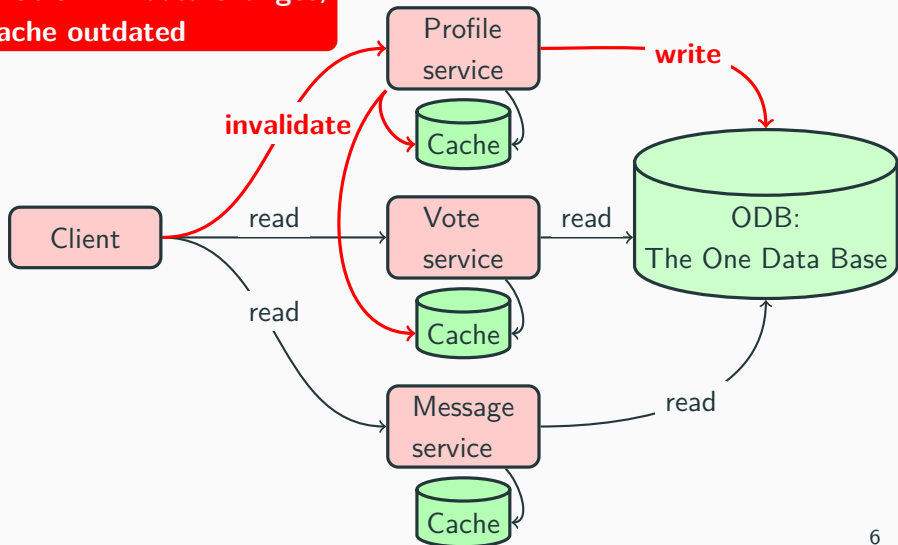
Reading data: “old school” approach

Problem: if data changes, cache outdated



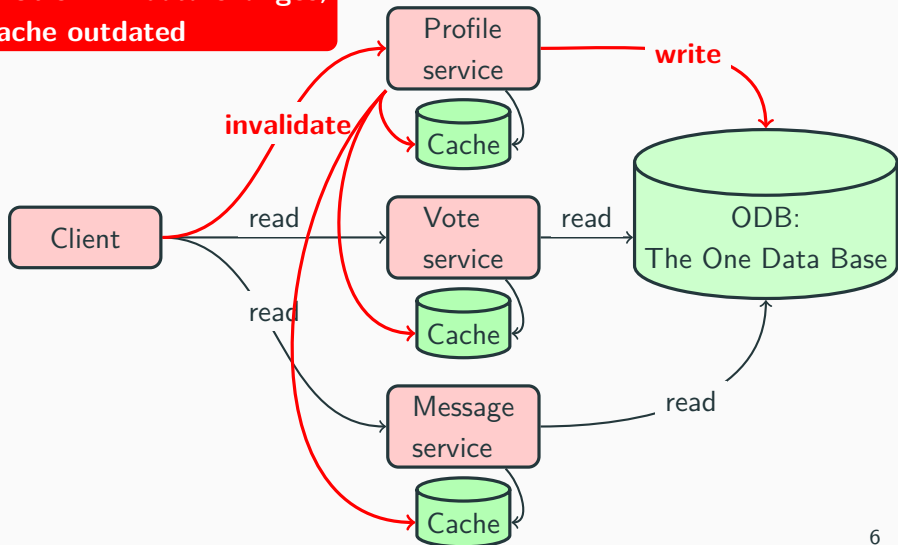
Reading data: “old school” approach

Problem: if data changes, cache outdated

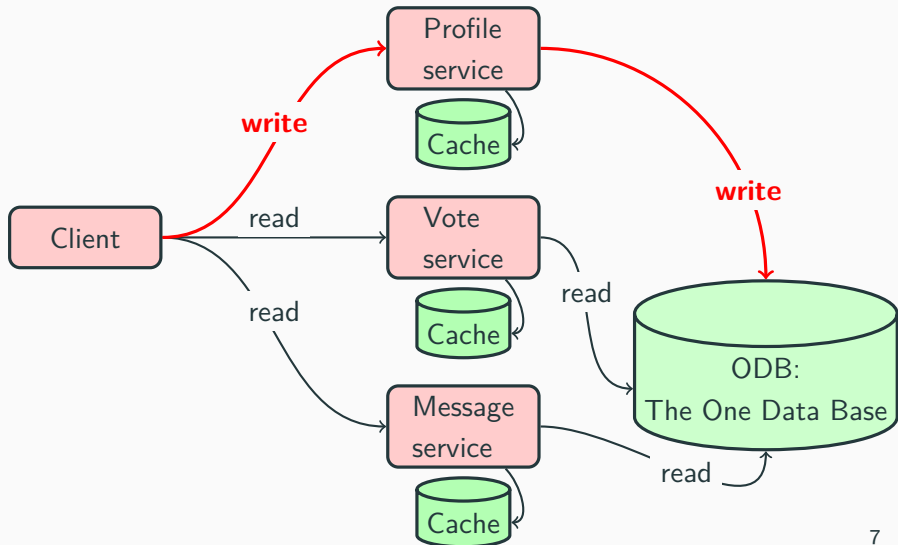


Reading data: “old school” approach

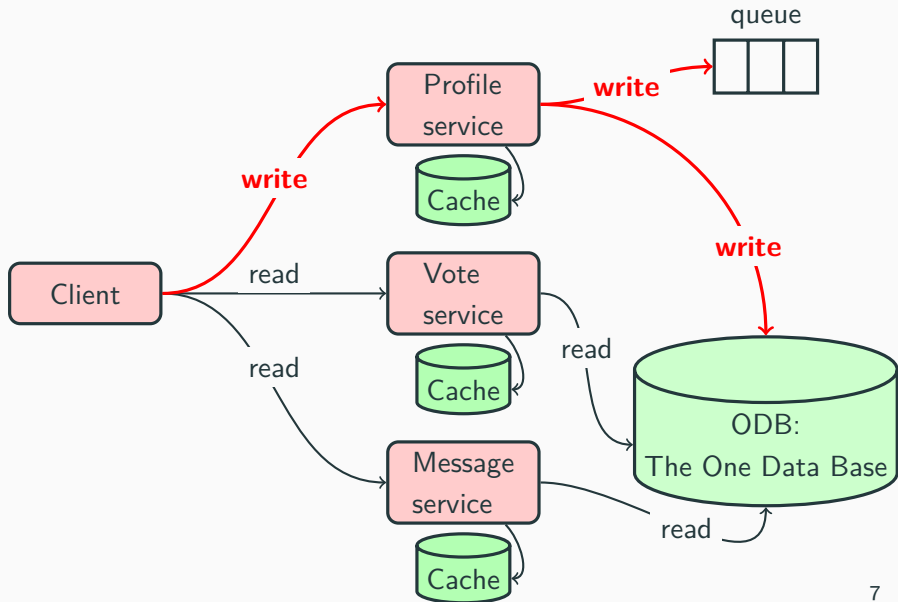
Problem: if data changes, cache outdated



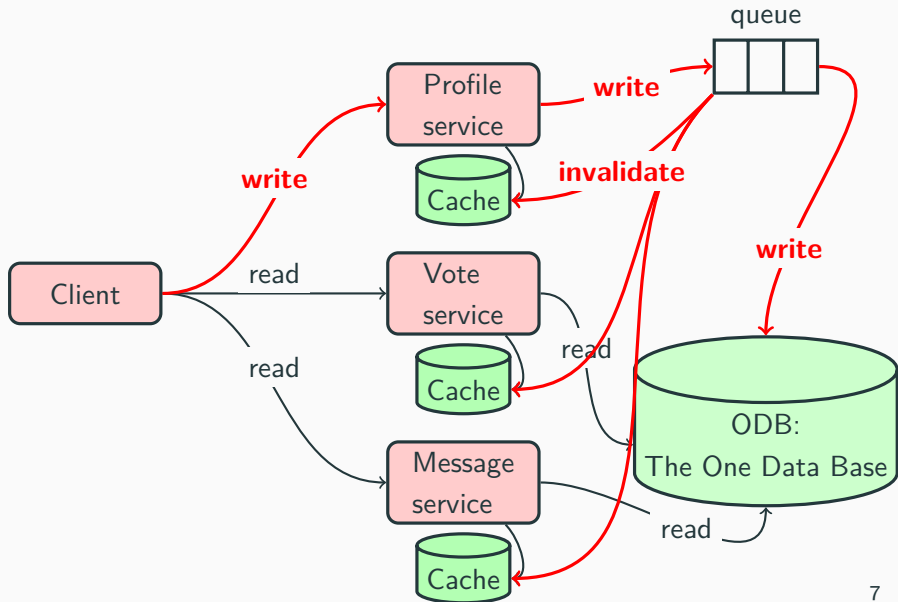
Reading data: decoupling microservices in space and time



Reading data: decoupling microservices in space and time

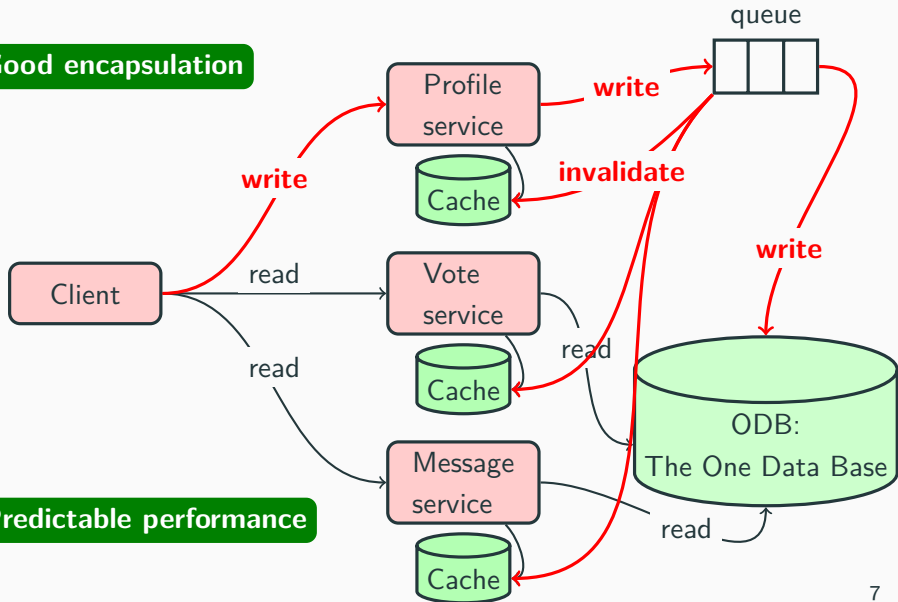


Reading data: decoupling microservices in space and time



Reading data: decoupling microservices in space and time

Good encapsulation



Predictable performance



Database abstraction

key	value

SET key1 = value1;

key	value

SET key1 = value1;

looks like

key	value
key1	value1

SET key1 = value1;

looks like

key	value
key1	value1

but actually

The Log



SET key1 = value1;

looks like

key	value
key1	value1

but actually



(0, key1, value1)

SET key1 = value1;

SET key2 = value2;

looks like

key	value
key1	value1
key2	value2

append



(0, key1, value1)

(1, key2, value2)

SET key1 = value1;

SET key2 = value2;

SET key1 = value3;

looks like

key	value
key1	value3
key2	value2

append



(0, key1, value1)

(1, key2, value2)

(2, key1, value3)

SET key1 = value1;
SET key2 = value2;
SET key1 = value3;

master

key	value
key1	value3
key2	value2

append



(0, key1, value1)
(1, key2, value2)
(2, key1, value3)

SET key1 = value1;
SET key2 = value2;
SET key1 = value3;

append



(0, key1, value1)
(1, key2, value2)
(2, key1, value3)

master

key	value
key1	value3
key2	value2

backup

key	value

SET key1 = value1;
SET key2 = value2;
SET key1 = value3;




master

key	value
key1	value3
key2	value2

backup

key	value
key1	value3
key2	value2



(0, key1, value1)
(1, key2, value2)
(2, key1, value3)



=



?



=



?

Log is really FIFO: order matters!



=




?


Log is really FIFO: order matters!

Log is permanent – Queue is transient


Stateful Stream Processing




Stateful Stream Processing




Stateful Stream Processing




Stateful Stream Processing



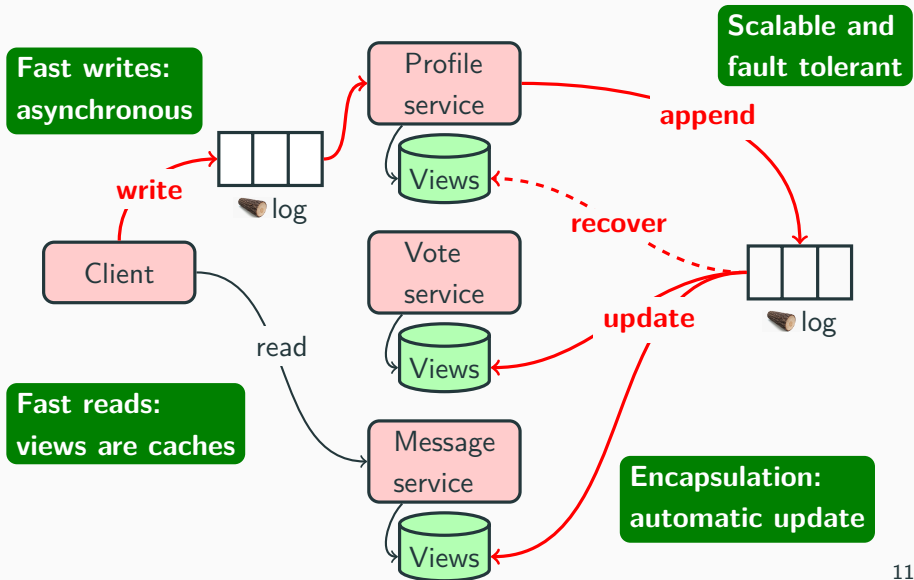
Stateful Stream Processing



Stateful Stream Processing

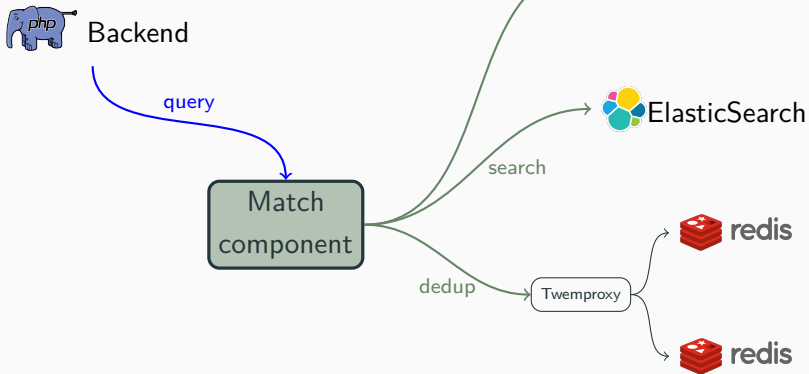



Stateful Stream Processing

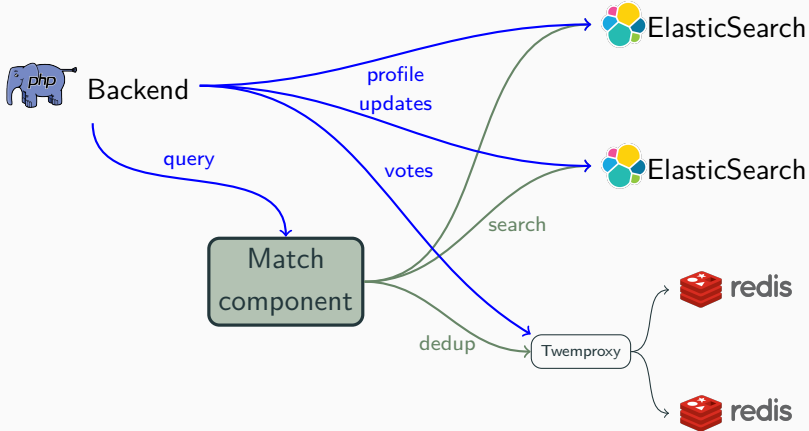


And now for something completely different. . .

A Real Example!









Backend

query

profile
updates

votes

search

dedup

Match
component

Twemproxy



ElasticSearch



ElasticSearch



redis



redis

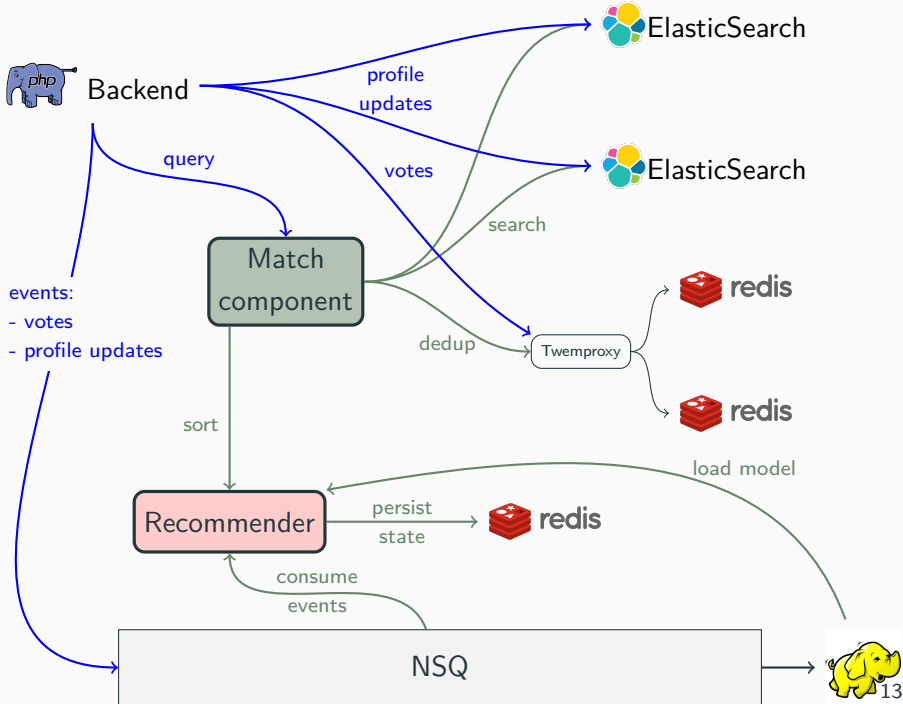
NSQ

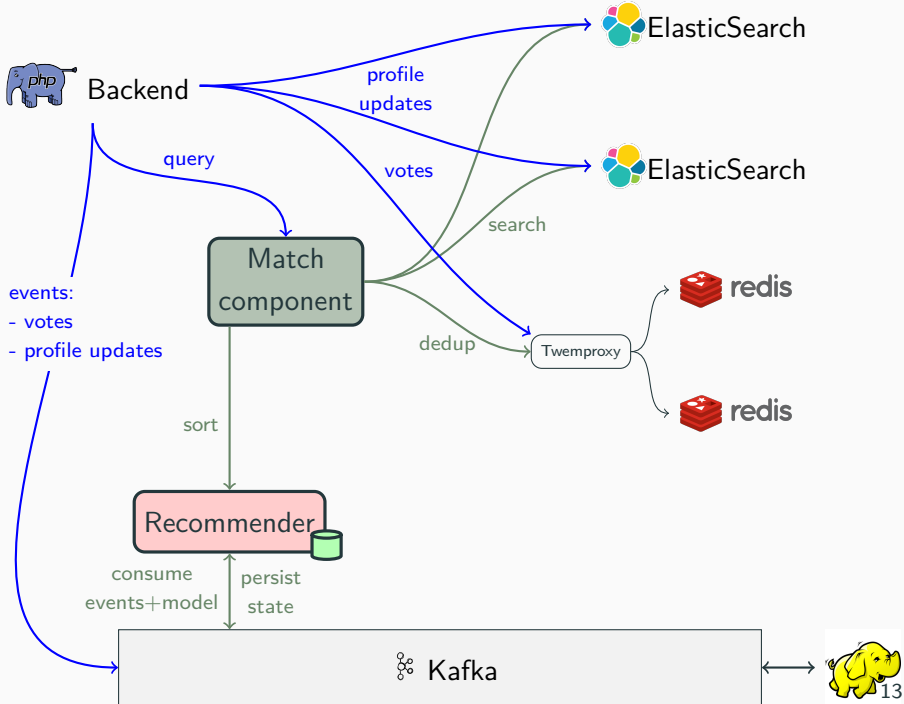


13

events:

- votes
- profile updates







Backend

query

profile
updates

votes

search

dedup

search

dedup

sort

consume
events+model

persist
state

consume
events

persist
state

consume
events

persist
state

Kafka



ElasticSearch



ElasticSearch



redis



redis

Twemproxy

Recommender

Edgeset

MatchSearch

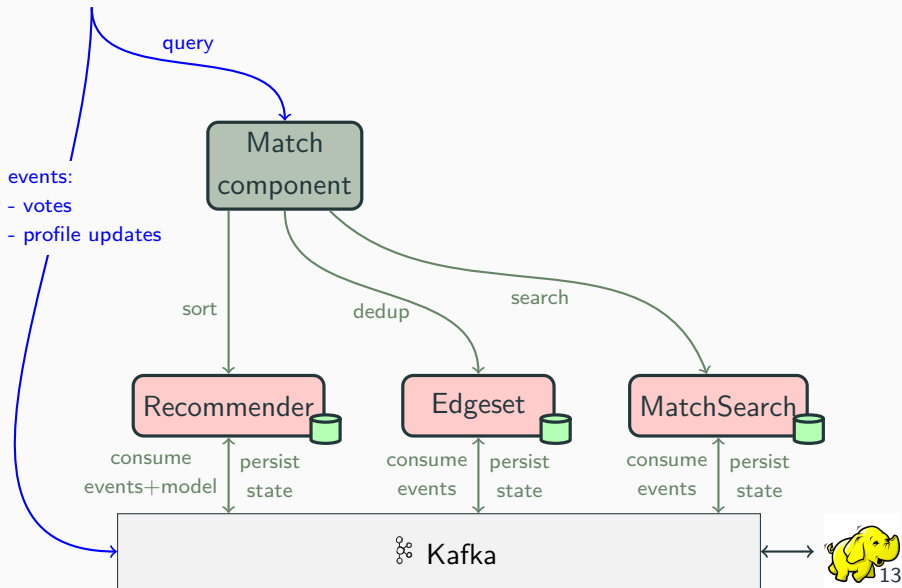
events:

- votes
- profile updates





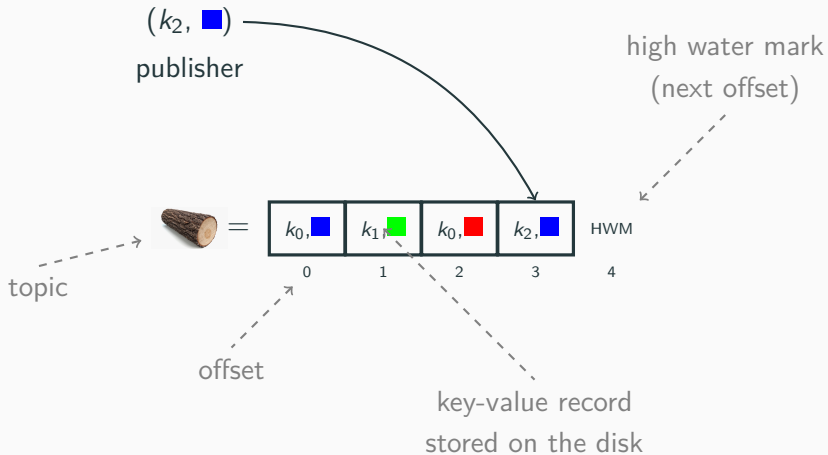
Backend




Kafka

topics, partitions, streams and tables


Kafka: Basics




Kafka: Basics




Kafka: Basics



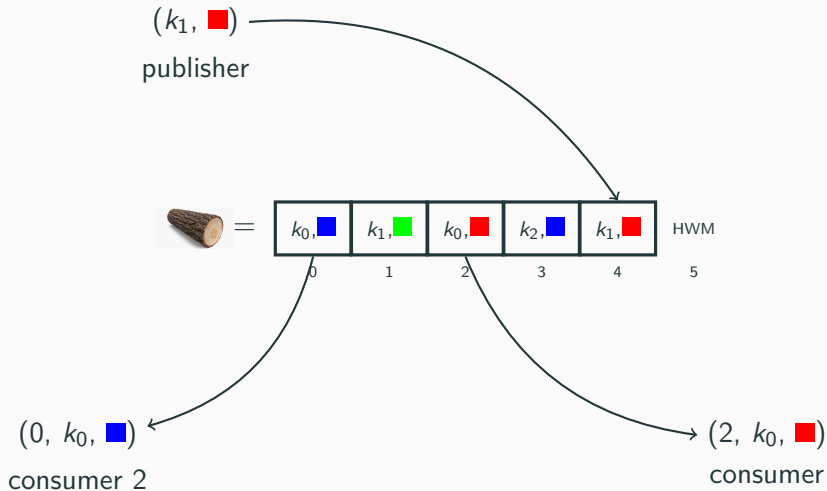
Kafka: Basics



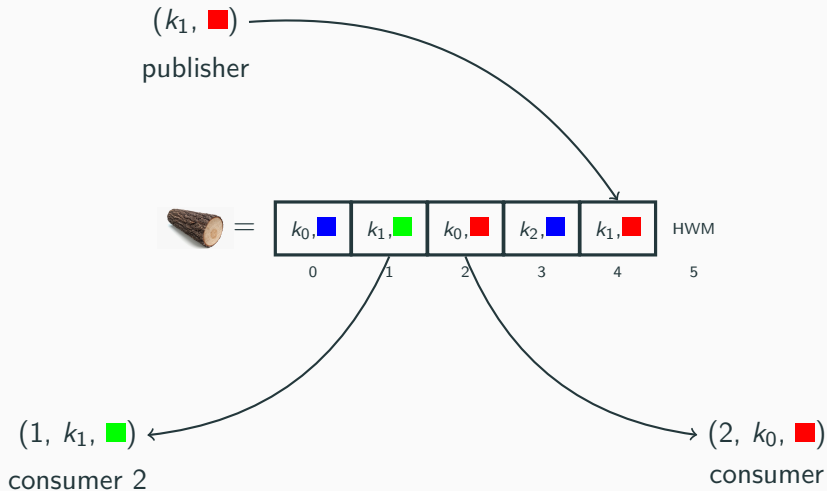
Kafka: Basics



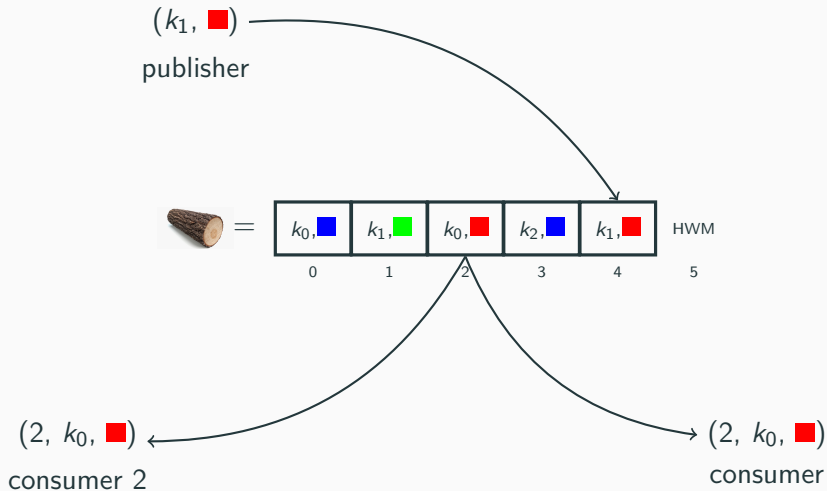
Kafka: Basics



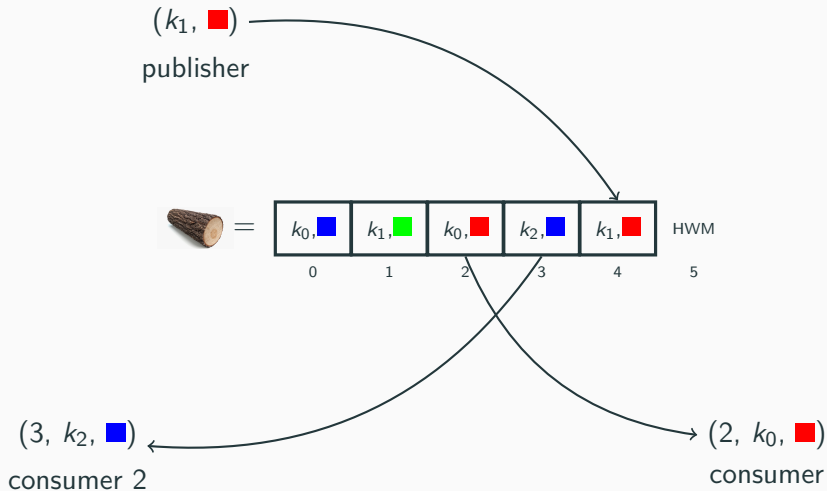
Kafka: Basics



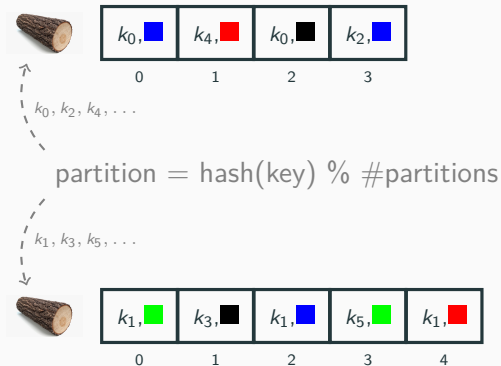
Kafka: Basics




Kafka: Basics





Kafka: Scalability and Fault Tolerance





Kafka: Scalability and Fault Tolerance





Kafka: Scalability and Fault Tolerance




Kafka: Scalability and Fault Tolerance




Kafka: Scalability and Fault Tolerance




Kafka: Scalability and Fault Tolerance




Kafka: Scalability and Fault Tolerance




Kafka: Scalability and Fault Tolerance




Kafka: Copartitioning



Kafka: Cleanup Policies



Stream



Table

Kafka: Cleanup Policies

delete after retention time
----->



Stream

compact, leaving latest value for each key
----->



Table

Kafka: Cleanup Policies

delete after retention time

----->



Stream

compact, leaving latest value for each key


----->



Table

Kafka: Cleanup Policies

delete after retention time
----->



Stream

compact, leaving latest value for each key
----->



Table

Kafka: Cleanup Policies

delete after retention time



4

5

Stream

compact, leaving latest value for each key



0

1

2

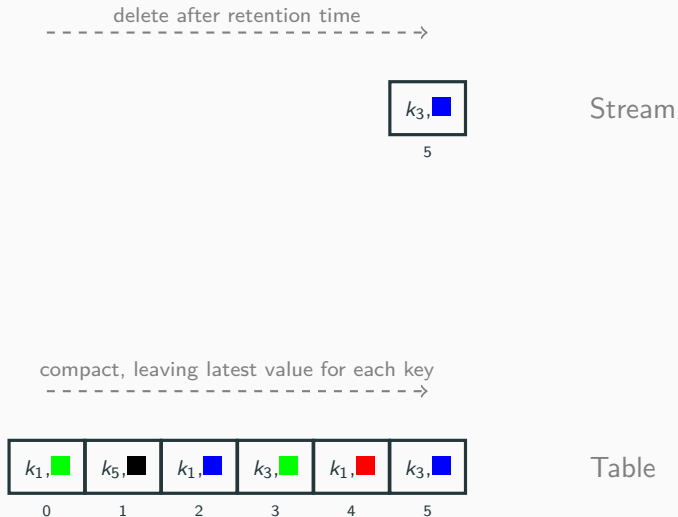
3

4


5

Table

Kafka: Cleanup Policies



Kafka: Cleanup Policies



Kafka: Cleanup Policies


delete after retention time
----->



5

Stream

compact, leaving latest value for each key
----->



1

2

3

4

5

Table

Kafka: Cleanup Policies

delete after retention time
----->



Stream

compact, leaving latest value for each key
----->



Table

Kafka: Cleanup Policies

delete after retention time
----->



5

Stream

compact, leaving latest value for each key
----->



1



4

5

Table

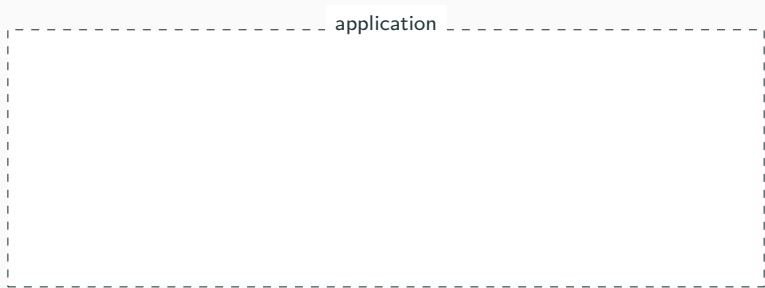
And now for something completely different. . .

Hands-on session!

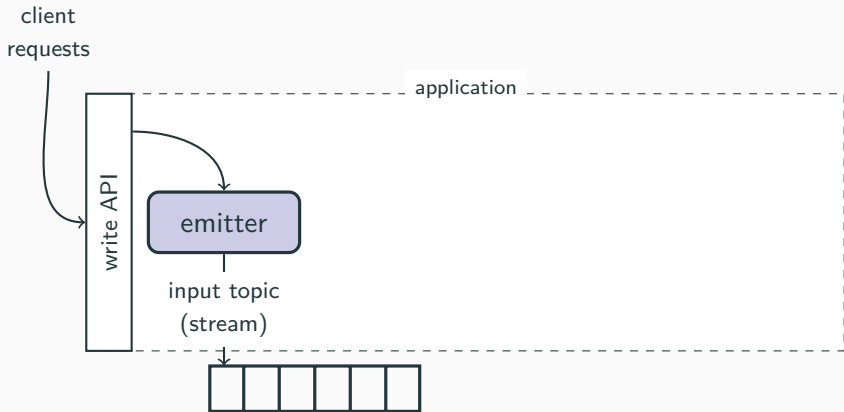
Learning Goka the “hard way”

components, code examples, and tasks

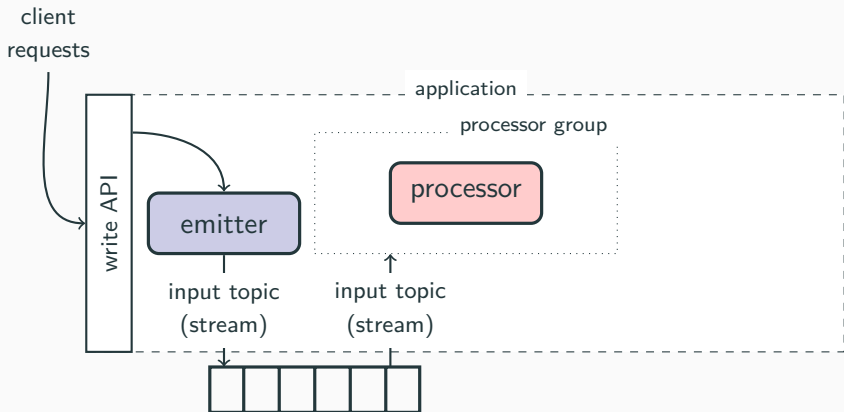
Goka Components



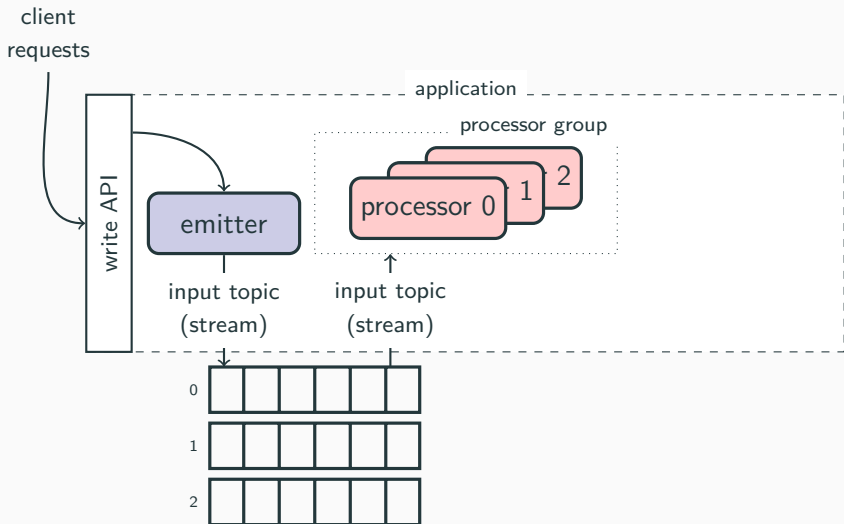
Goka Components



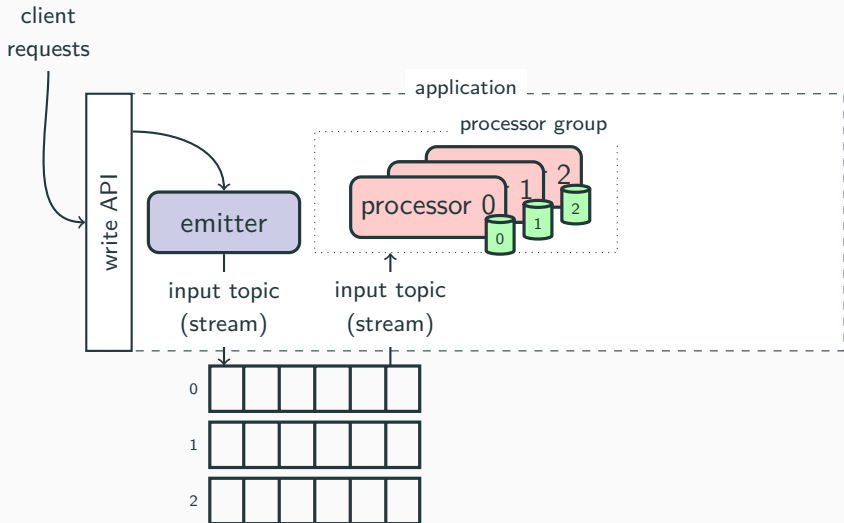
Goka Components



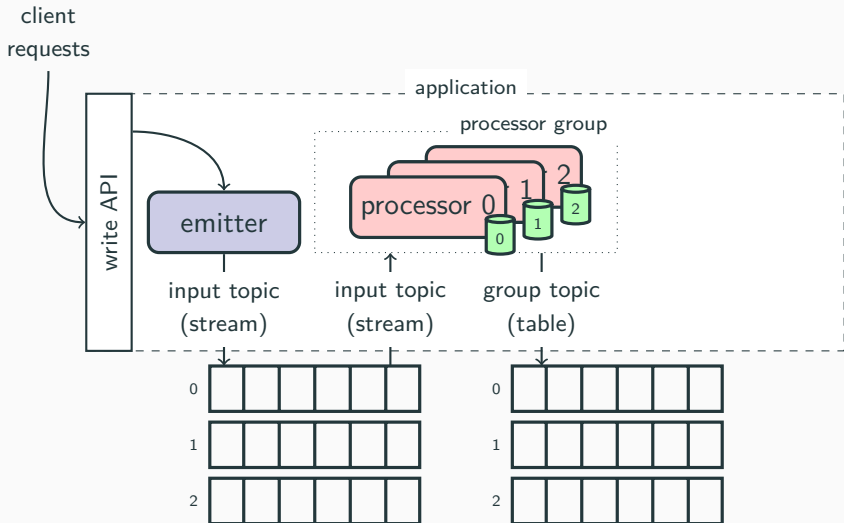
Goka Components



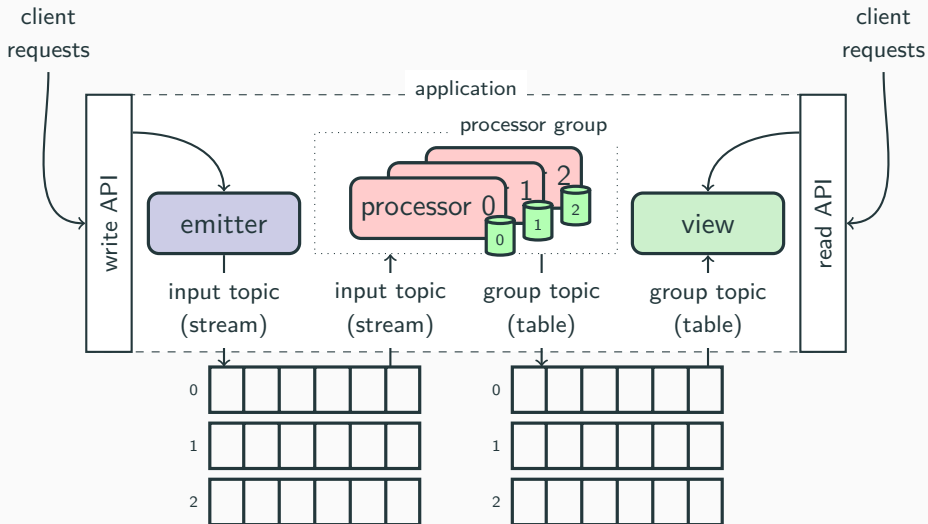
Goka Components



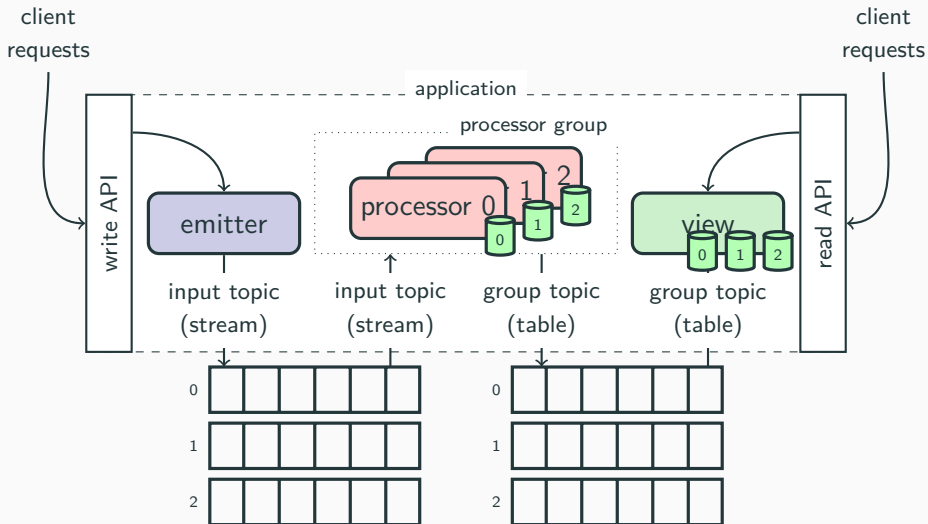
Goka Components



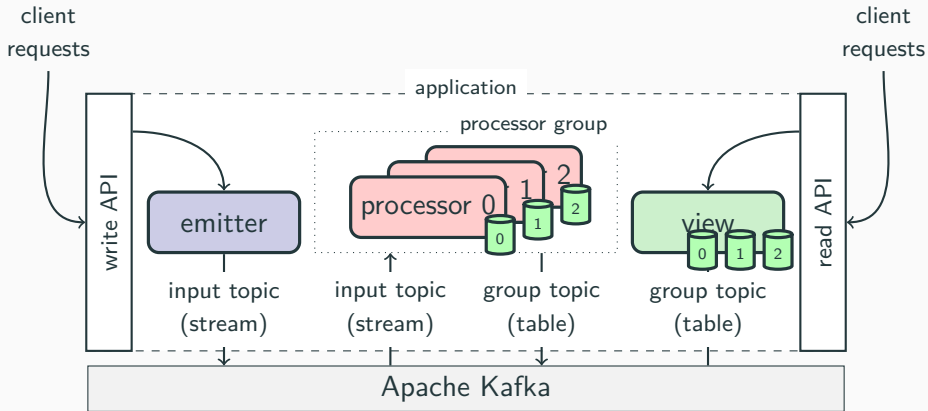
Goka Components




Goka Components



Goka Components




Emitter



```
// create new emitter
e, _ := goka.NewEmitter(
    brokers, // kafka:9092
    "input-stream", // target topic
    streamCodec, // message encoder
)

func handleRequest(req Request) {
    // emit asynchronously
    e.Emit(req.Key, req.Content)
}
```

View



```
// create view
v, _ := goka.NewView(
    brokers,
    "mygroup-table",
    tableCodec,
)

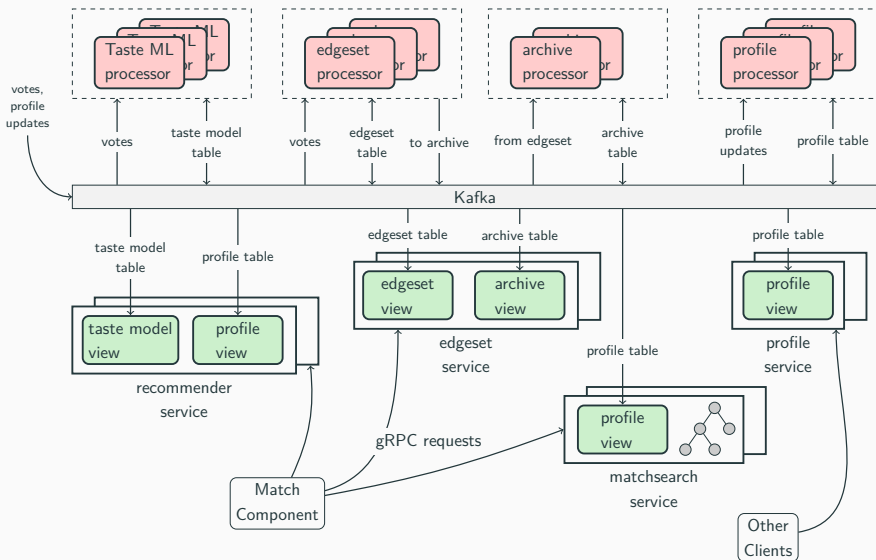
// start view
go v.Run(context.Background())

func handleRequest(req Request) {
    // access view as a kv table
    data, _ := v.Get(req.Key)
    if data == nil {
        // not found
    }
    // safe to cast
    value := data.(*tableEntry)
}
```

Live coding **and** live drawing. . .

And now for something somewhat different. . .

A Detailed Real Example!



Closing Discussion

you choose: tips, monitoring, testing, . . .

Take aways...

- Streaming microservices are cool!
- Kafka + Golang make a great combo!
- If you need state, consider Goka – it's awesome!

Where to go from here?

- **Try some examples**

- `github.com/lovoo/goka/tree/master/examples`

- `github.com/lovoo/cofire`

- **Read the 2-page wiki**

- `github.com/lovoo/goka/wiki/Introduction`

- `github.com/lovoo/goka/wiki/Tips`

- **Questions, issues, PRs**

- `github.com/lovoo/goka/issues`

More about stream processing...

- **Readings**

- I Heart Logs – Jay Kreps

- Making sense of stream processing – Martin Kleppmann

- **Video**

- Turning the database inside out with Samza – Kleppmann

- <https://bit.ly/2b7zjsx>

- **Moath's presentation tomorrow:**

- Using Apache Kafka with Golang

Feedback? Want to talk? Get in touch!

Diogo – `diogo.behrens@volkswagen.de`

Franz – `franz.eichhorn@lovoo.com`